

Kubernetes Beginners Workshop

@Vendevio by Tobias Mühlberger

26.05.2021

Workshop Structure

- Intro
- Terminology
- Architecture and Concepts
- Hands-on Exercises

Materials: github.com/t-muehlberger/k8s-beginners-workshop-vendevo



Prerequisites

- Basic Docker/Container knowledge
- Installed **kubect1** locally
- Access to a Kubernetes cluster
 - minikube
 - microk8s
 - **Vendevio Dev Cluster**



What you will learn

- How to interact with Kubernetes using **kubectl** and the rancher web-ui
- How to deploy an application to Kubernetes
- How to troubleshoot problems with your application
 - Access logs of your application
 - Access ports of your application
 - Access your database
- Basic insights into what happens under the hood



What is Kubernetes

- Container orchestration platform
- for automating deployment, scaling, and management of containerized applications
- across clusters of hosts.
- “The Operating System of the Cloud”



Why Kubernetes

- Infrastructure as Code
- Desired state configuration
- Self-healing
- Cross-Cloud portability
- Open source ecosystem
- Layer between Devs and Ops



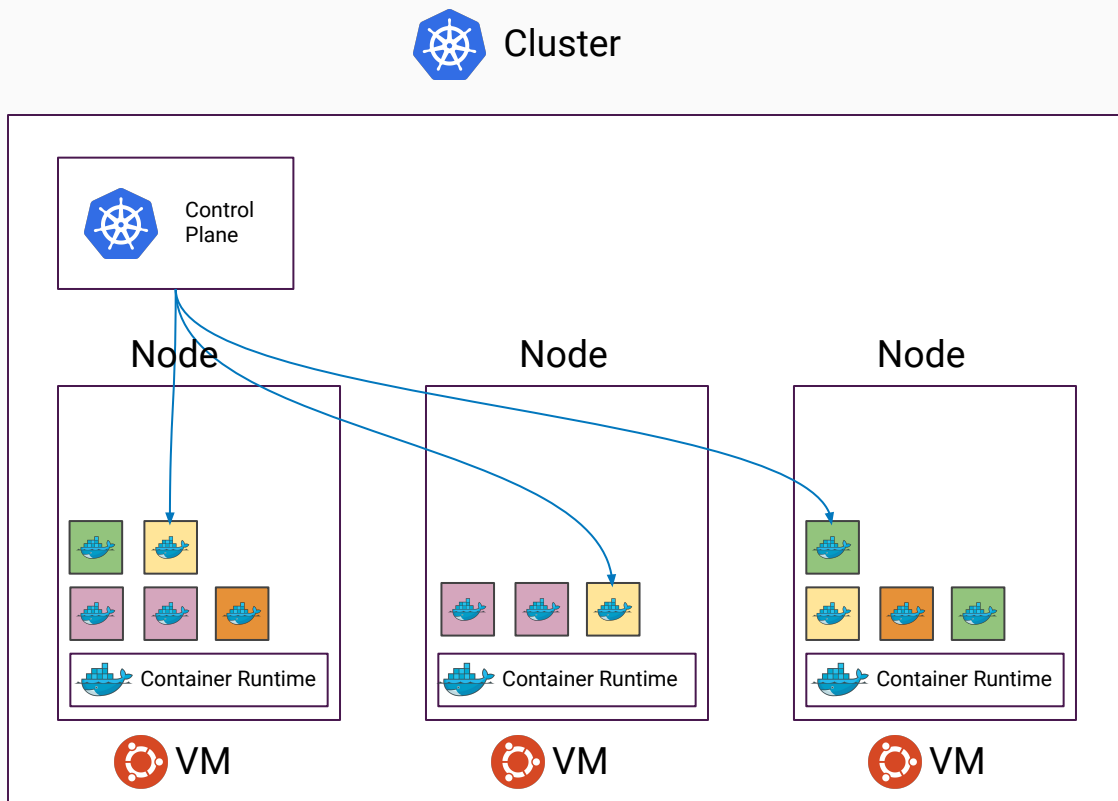
Kubernetes vs. Docker vs. Docker-Compose

- Declarative vs. Imperative
- Multiple Servers vs. Single Server



Terminology

- Cluster
- Node
- Pod
- Workload
- Control Plane
- Container runtime
- Kubernetes Objects
- Helm Charts



Cheat Sheet

```
kubectl get all -n <namespace>
```

```
kubectl get <kind> <name>
```

```
kubectl get service postgres
```

```
kubectl describe <kind>/<name>
```

```
kubectl logs -f <pod-name>
```

```
kubectl apply -f config.yml
```

```
kubectl delete -f config.yml
```

Exercise 1: Namespace

- “Kubernetes namespaces help different projects, teams, or customers to share a Kubernetes cluster.”
- Namespaces help organize and avoid naming conflicts
- Separate projects should be deployed in separate namespaces

```
kubectl create namespace <first_name>
```



Exercise 2: Deployment

- Describes a Workload
- Deployment Controller will create Pods
- Controls Replication
- Rolling Updates
- Use Deployments rather than creating a Pod manually



Exercise 3: Service

- Abstraction for access to Pod
- Service Discovery
 - Caller does not need to know the name or IP address of the Pod
 - A single DNS name for a set of Pods
 - Acts as internal load balancer for replicated Pods
- Type
 - **ClusterIP**
 - NodePort
 - ...



Exercise 4: Ingress

- Reverse Proxy
 - Nginx
- Alternative to Traefik in Docker-Compose Deployment
- Routing
- SSL Termination



Exercise 5: Persistence

- Required for stateful applications such as databases
- Persistent Volume Claim (PVC) is an Abstraction
- Storage Controller will create a Persistent Volume (PV) according to specification
- Different Storage Classes for different type of storage



Exercise 5.1: Postgres

- Connect to Postgres via Port-Forwarding
 - `kubectl port-forward -n <namespace> service/postgres 5432:5432`
- Access to Console
 - `kubectl exec -ti -n <namespace> <postgres-pod> -- /bin/bash`
- Create DB Dump
 - `kubectl exec -n <namespace> <postgres-pod> -- pg_dump -U postgres > db-dump.sql`
- Restore DB Dump
 - `kubectl exec -ti -n <namespace> <postgres-pod> -- psql -U postgres postgres < db-dump.sql`



Exercise 6: ConfigMaps + Secrets

- Enables separation between configuration and Deployment
- Config can be used as
 - Environment Variable
 - Mounted config File
- Protect Secrets



Exercise 7: (Cron) Jobs

- One time Job or recurring Job
- Should be used for:
 - DB Migrations
 - Backups



Extending Kubernetes

- Custom Resource Definition (CRD)
- Operators
- Example
 - Velero Backup CRD



Useful Ressources

<https://kubernetes.io/docs/concepts/>

<https://kubernetes.io/docs/tutorials/>

<https://kubernetes.io/docs/setup/best-practices/>

<https://www.youtube.com/watch?v=X48VuDVv0do>

